

Insurance Claims Fraud Detection


```
In [1]: # Importing Necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('insurance_claim_fraud.csv')
df
```

```
Out[2]:
```

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_
0	328	48	521585	17-10-2014	OH	250/500	
1	228	42	342868	27-06-2006	IN	250/500	
2	134	29	687698	06-09-2000	OH	100/300	
3	256	41	227811	25-05-1990	IL	250/500	
4	228	44	367455	06-06-2014	IL	500/1000	
...	
995	3	38	941851	16-07-1991	OH	500/1000	
996	285	41	186934	05-01-2014	IL	100/300	
997	130	34	918516	17-02-2003	OH	250/500	
998	458	62	533940	18-11-2011	IL	500/1000	
999	456	60	556080	11-11-1996	OH	250/500	

1000 rows × 40 columns



So here we have 1000 rows and 40 columns in data

- fraud_reported is our target variable

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   months_as_customer                    1000 non-null   int64
1   age                                   1000 non-null   int64
2   policy_number                         1000 non-null   int64
3   policy_bind_date                     1000 non-null   object
4   policy_state                         1000 non-null   object
5   policy_csl                           1000 non-null   object
6   policy_deductable                    1000 non-null   int64
7   policy_annual_premium                 1000 non-null   float64
8   umbrella_limit                       1000 non-null   int64
9   insured_zip                          1000 non-null   int64
10  insured_sex                          1000 non-null   object
11  insured_education_level               1000 non-null   object
12  insured_occupation                   1000 non-null   object
13  insured_hobbies                       1000 non-null   object
14  insured_relationship                 1000 non-null   object
15  capital-gains                        1000 non-null   int64
16  capital-loss                         1000 non-null   int64
17  incident_date                        1000 non-null   object
18  incident_type                        1000 non-null   object
19  collision_type                       1000 non-null   object
20  incident_severity                    1000 non-null   object
21  authorities_contacted                1000 non-null   object
22  incident_state                       1000 non-null   object
23  incident_city                        1000 non-null   object
24  incident_location                    1000 non-null   object
25  incident_hour_of_the_day              1000 non-null   int64
26  number_of_vehicles_involved           1000 non-null   int64
27  property_damage                      1000 non-null   object
28  bodily_injuries                      1000 non-null   int64
29  witnesses                            1000 non-null   int64
30  police_report_available               1000 non-null   object
31  total_claim_amount                   1000 non-null   int64
32  injury_claim                         1000 non-null   int64
33  property_claim                       1000 non-null   int64
34  vehicle_claim                        1000 non-null   int64
35  auto_make                            1000 non-null   object
36  auto_model                           1000 non-null   object
37  auto_year                            1000 non-null   int64
38  fraud_reported                       1000 non-null   object
39  _c39                                0 non-null      float64
dtypes: float64(2), int64(17), object(21)
memory usage: 312.6+ KB
```

- Here we have 19 numerical and 21 object columns present in the dataset.

```
In [4]: df.isnull().sum()
```

```
Out[4]: months_as_customer    0
age                           0
policy_number                 0
policy_bind_date              0
policy_state                  0
policy_csl                    0
policy_deductable             0
policy_annual_premium         0
umbrella_limit                0
insured_zip                   0
insured_sex                   0
insured_education_level       0
insured_occupation            0
insured_hobbies               0
insured_relationship          0
capital-gains                 0
capital-loss                  0
incident_date                 0
incident_type                 0
collision_type                0
incident_severity             0
authorities_contacted         0
incident_state                0
incident_city                 0
incident_location             0
incident_hour_of_the_day      0
number_of_vehicles_involved   0
property_damage               0
bodily_injuries               0
witnesses                     0
police_report_available       0
total_claim_amount            0
injury_claim                  0
property_claim                0
vehicle_claim                 0
auto_make                     0
auto_model                    0
auto_year                     0
fraud_reported                0
_c39                          1000
dtype: int64
```

- There is no null value present in the columns except _c39 column so dropping it

```
In [5]: # dropping _c39 column so dropping it.
df.drop('_c39',axis=1,inplace=True)
```

```
In [6]: df.shape
```

```
Out[6]: (1000, 39)
```

```
In [7]: # printing all data type and their unique values
for column in df.columns:
    if df[column].dtype==object:
        print(df[column].value_counts())
        print('*****')
```

```
05-08-1992    3
01-01-2006    3
28-04-1992    3
29-01-1998    2
21-12-2002    2
..
25-11-1994    1
12-02-2009    1
28-03-2001    1
17-12-2003    1
10-12-2014    1
Name: policy_bind_date, Length: 951, dtype: int64
*****
OH      352
IL      338
IN      310
Name: policy_state, dtype: int64
*****
250/500      351
100/200      310
```

```
In [8]: # replacing '?' with No Info
df=df.replace('?', 'No Info')
```

- Some of the rows having ? sign replacing them with No Info

```
In [9]: # printing all data type and their unique values
for column in df.columns:
    if df[column].dtype==object:
        print(df[column].value_counts())
        print('*****')
```

05-08-1992	3
01-01-2006	3
28-04-1992	3
29-01-1998	2
21-12-2002	2
..	
25-11-1994	1
12-02-2009	1
28-03-2001	1
17-12-2003	1
10-12-2014	1

Name: policy_bind_date, Length: 951, dtype: int64

OH	352
IL	338
IN	310

Name: policy_state, dtype: int64

250/500	351
100/300	349
500/1000	300

Name: policy_csl, dtype: int64

FEMALE	537
MALE	463

Name: insured_sex, dtype: int64

JD	161
High School	160
Associate	145
MD	144
Masters	143
PhD	125
College	122

Name: insured_education_level, dtype: int64

machine-op-inspct	93
prof-specialty	85
tech-support	78
sales	76
exec-managerial	76
craft-repair	74
transport-moving	72
other-service	71
priv-house-serv	71
armed-forces	69
adm-clerical	65
protective-serv	63
handlers-cleaners	54
farming-fishing	53

Name: insured_occupation, dtype: int64

reading	64
paintball	57
exercise	57
bungie-jumping	56
camping	55

golf	55
movies	55
kayaking	54
yachting	53
hiking	52
video-games	50
base-jumping	49
skydiving	49
board-games	48
polo	47
chess	46
dancing	43
sleeping	41
cross-fit	35
basketball	34

Name: insured_hobbies, dtype: int64

own-child	183
other-relative	177
not-in-family	174
husband	170
wife	155
unmarried	141

Name: insured_relationship, dtype: int64

02-02-2015	28
17-02-2015	26
07-01-2015	25
24-01-2015	24
10-01-2015	24
04-02-2015	24
19-01-2015	23
08-01-2015	22
30-01-2015	21
13-01-2015	21
12-02-2015	20
31-01-2015	20
22-02-2015	20
06-02-2015	20
21-01-2015	19
14-01-2015	19
21-02-2015	19
01-01-2015	19
12-01-2015	19
23-02-2015	19
14-02-2015	18
20-01-2015	18
18-01-2015	18
03-01-2015	18
01-02-2015	18
28-02-2015	18
25-02-2015	18
09-01-2015	17
24-02-2015	17
06-01-2015	17
08-02-2015	17
26-02-2015	17

15-02-2015	16
16-02-2015	16
05-02-2015	16
16-01-2015	16
13-02-2015	16
15-01-2015	15
28-01-2015	15
17-01-2015	15
18-02-2015	15
27-02-2015	14
22-01-2015	14
20-02-2015	14
23-01-2015	13
09-02-2015	13
03-02-2015	13
27-01-2015	13
01-03-2015	12
04-01-2015	12
26-01-2015	11
29-01-2015	11
02-01-2015	11
07-02-2015	10
10-02-2015	10
11-02-2015	10
19-02-2015	10
25-01-2015	10
11-01-2015	9
05-01-2015	7

Name: incident_date, dtype: int64

Multi-vehicle Collision	419
Single Vehicle Collision	403
Vehicle Theft	94
Parked Car	84

Name: incident_type, dtype: int64

Rear Collision	292
Side Collision	276
Front Collision	254
No Info	178

Name: collision_type, dtype: int64

Minor Damage	354
Total Loss	280
Major Damage	276
Trivial Damage	90

Name: incident_severity, dtype: int64

Police	292
Fire	223
Other	198
Ambulance	196
None	91

Name: authorities_contacted, dtype: int64

NY	262
SC	248


```

WV      217
NC      110
VA      110
PA       30
OH       23
Name: incident_state, dtype: int64
*****

Springfield    157
Arlington      152
Columbus       149
Northbend      145
Hillsdale      141
Riverwood      134
Northbrook     122
Name: incident_city, dtype: int64
*****

8983 Tree St          1
1589 Best Ave         1
3847 Elm Hwy          1
4780 Best Drive       1
8689 Maple Hwy        1
..
5783 Oak Ave          1
3790 Andromedia Hwy   1
1507 Solo Ave         1
4629 Elm Ridge        1
6484 Tree Drive       1
Name: incident_location, Length: 1000, dtype: int64
*****

No Info    360
NO         338
YES        302
Name: property_damage, dtype: int64
*****

No Info    343
NO         343
YES        314
Name: police_report_available, dtype: int64
*****

Suburu      80
Saab        80
Dodge       80
Nissan       78
Chevrolet   76
Ford        72
BMW         72
Toyota      70
Audi        69
Accura      68
Volkswagen  68
Jeep        67
Mercedes    65
Honda       55
Name: auto_make, dtype: int64
*****

RAM          43
Wrangler     42

```

A3	37
Neon	37
MDX	36
Jetta	35
Passat	33
Legacy	32
A5	32
Pathfinder	31
Malibu	30
Forrester	28
92x	28
Camry	28
F150	27
95	27
E400	27
93	25
Grand Cherokee	25
Escape	24
Maxima	24
Tahoe	24
Ultima	23
X5	23
Highlander	22
Silverado	22
Civic	22
Fusion	21
TL	20
Corolla	20
CRV	20
Impreza	20
ML350	20
3 Series	18
C300	18
X6	16
M5	15
Accord	13
RSX	12

Name: auto_model, dtype: int64

N 753

Y 247

Name: fraud_reported, dtype: int64

Description of Dataset

```
In [10]: # statistcla summary of numerical columns
df.describe()
```

```
Out[10]:
```

	months_as_customer	age	policy_number	policy_deductable	policy_annual_premi
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.0000
mean	203.954000	38.948000	546238.648000	1136.000000	1256.4061
std	115.113174	9.140287	257063.005276	611.864673	244.1673
min	0.000000	19.000000	100804.000000	500.000000	433.3300
25%	115.750000	32.000000	335980.250000	500.000000	1089.6075
50%	199.500000	38.000000	533135.000000	1000.000000	1257.2000
75%	276.250000	44.000000	759099.750000	2000.000000	1415.6950
max	479.000000	64.000000	999435.000000	2000.000000	2047.5900

This gives the statistical information of the numerical columns . The summary of this dataset looks perfect

From the above description we can observe the following things.

- The count of all the columns are same which means there are no missing values in the dataset.
- The mean value is greater than the median(50%) in most of the columns which means the data is skewed to right in these columns
- The data in the few columns have mean value less than median that means the data is skewed to left
- By summerising the data we can observe there is huge difference between 75% and max in most of the columns hence there are outliers present in the data which we will remove them leter on using appropriate methods.
- We can also notice the Standard deviation , min, 25% percentile values from this describe method.

```
In [11]: df['umbrella_limit'].value_counts()
```

```
Out[11]:
```

0	798
6000000	57
5000000	46
4000000	39
7000000	29
3000000	12
8000000	8
9000000	5
2000000	3
10000000	2
-1000000	1

Name: umbrella_limit, dtype: int64


- We can see there is one row in negative value must be the mistake or not sure so dropping the row

In [12]: `df.loc[df['umbrella_limit']== -1000000]`

Out[12]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_
290	284	42	526039	04-05-1995	OH	100/300	

1 rows × 39 columns



In [13]: `# dropping that value row(one row only)`
`df.drop(df[df['umbrella_limit']== -1000000].index,inplace=True)`

In [14]: `# dividing ['incident_data'] into three other columns`
`df['incident_Date']=df['incident_date'].str.split('-').str[0]`
`df['incident_Month']=df['incident_date'].str.split('-').str[1]`
`df['incident_Year']=df['incident_date'].str.split('-').str[2]`

In [15]: `# in incident data all the incident is from 2015 so dropping the incident_year`
`df['incident_Year'].value_counts()`
`df.drop('incident_Year',axis=1,inplace=True)`

In [16]: `# dividing policy_bind_date into three other columns`
`df['policy_bind_Date']=df['policy_bind_date'].str.split('-').str[0]`
`df['policy_bind_Month']=df['policy_bind_date'].str.split('-').str[1]`
`df['policy_bind_Year']=df['policy_bind_date'].str.split('-').str[2]`


`# dropping policy_bind_date column`
`df.drop('policy_bind_date',axis=1,inplace=True)`

In [17]: `df.head()`

Out[17]:

	months_as_customer	age	policy_number	policy_state	policy_csl	policy_deductable	policy_a
0	328	48	521585	OH	250/500	1000	
1	228	42	342868	IN	250/500	2000	
2	134	29	687698	OH	100/300	2000	
3	256	41	227811	IL	250/500	2000	
4	228	44	367455	IL	500/1000	1000	

5 rows × 43 columns



```
In [18]: # dropping the policy_number and other useless columns as well
df.drop(['policy_number', 'insured_zip', 'incident_location'], axis=1, inplace=True)
```

```
In [19]: df.head()
```

```
Out[19]:
```

	months_as_customer	age	policy_state	policy_csl	policy_deductable	policy_annual_premium
0	328	48	OH	250/500	1000	1406.91
1	228	42	IN	250/500	2000	1197.22
2	134	29	OH	100/300	2000	1413.14
3	256	41	IL	250/500	2000	1415.74
4	228	44	IL	500/1000	1000	1583.91

5 rows × 40 columns



```
In [20]: df.shape
```


```
Out[20]: (999, 40)
```

```
In [21]: df['fraud_reported'].value_counts()
```

```
Out[21]: N    752
         Y    247
         Name: fraud_reported, dtype: int64
```

```
In [22]: sns.countplot(df['fraud_reported'])
```

```
Out[22]: <AxesSubplot:xlabel='fraud_reported', ylabel='count'>
```



- We can notice data is imbalance we have to deal with it.

EDA

```
In [23]: df.head()
```

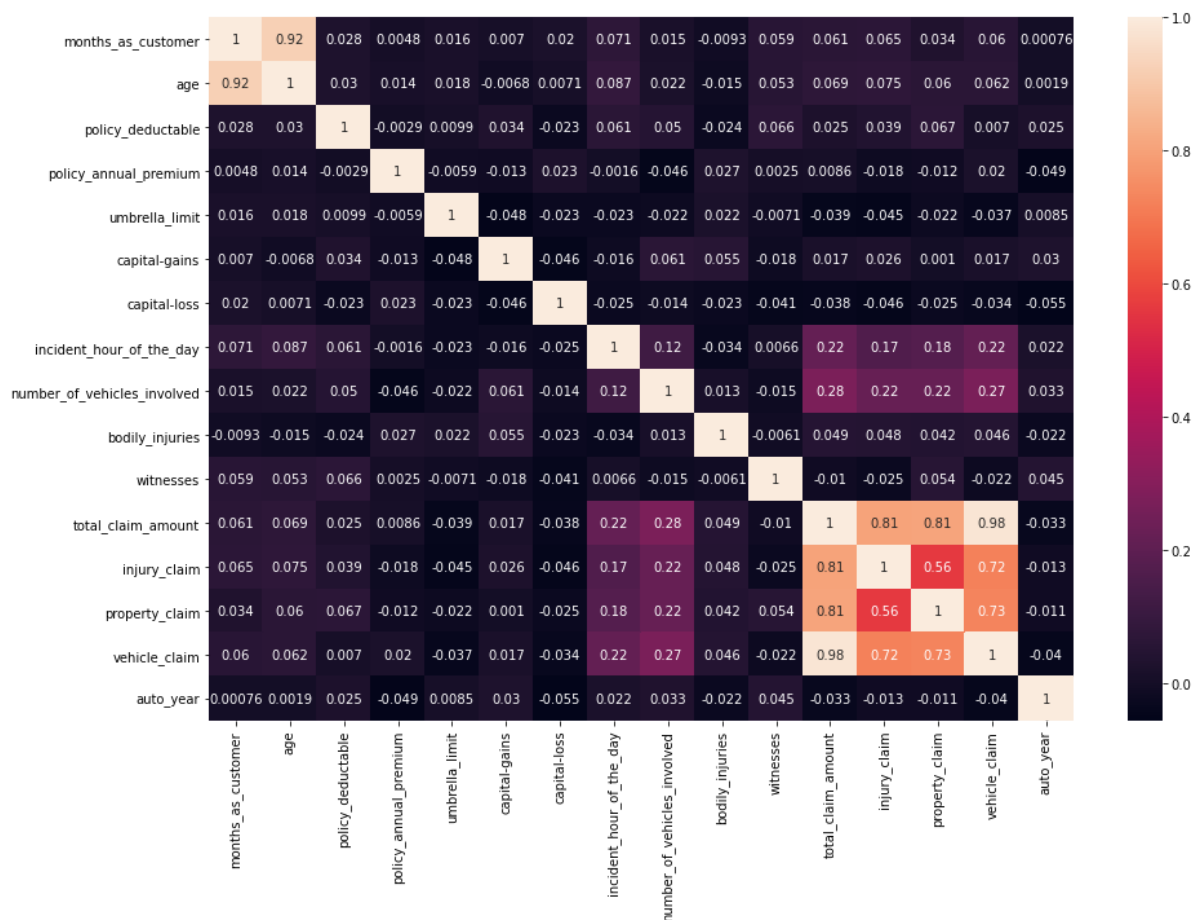
```
Out[23]:
```

	months_as_customer	age	policy_state	policy_csl	policy_deductable	policy_annual_premium
0	328	48	OH	250/500	1000	1406.91
1	228	42	IN	250/500	2000	1197.22
2	134	29	OH	100/300	2000	1413.14
3	256	41	IL	250/500	2000	1415.74
4	228	44	IL	500/1000	1000	1583.91

5 rows × 40 columns

```
In [24]: plt.figure(figsize=(15,10))
sns.heatmap(df.corr(),annot=True)
```


```
Out[24]: <AxesSubplot:>
```



- We can notice total_claim_amount, injury_claim, property_claim and vehicle_claim are highly correlated with each other .
- All other columns are very least correlated with each other.

```
In [25]: plt.figure(figsize=(15,7))
sns.violinplot(y='total_claim_amount',x='fraud_reported',data=df)
```


```
Out[25]: <AxesSubplot:xlabel='fraud_reported', ylabel='total_claim_amount'>
```



- Most of fraud reported found where total claimed amount 50000 to 70000

```
In [26]: plt.figure(figsize=(15,7))
sns.countplot(x='number_of_vehicles_involved',hue='fraud_reported',data=df)
```


```
Out[26]: <AxesSubplot:xlabel='number_of_vehicles_involved', ylabel='count'>
```



- Most of the fraud report were found where number_of_vehicles_involved more than 1

```
In [27]: plt.figure(figsize=(15,7))
sns.countplot(x='incident_city',hue='fraud_reported',data=df)
```


```
Out[27]: <AxesSubplot:xlabel='incident_city', ylabel='count'>
```



- In all cities fraud reported counts are almost same.

```
In [28]: plt.figure(figsize=(15,7))
sns.countplot(x='incident_state',hue='fraud_reported',data=df)
```


```
Out[28]: <AxesSubplot:xlabel='incident_state', ylabel='count'>
```



- Most of the fraud reported were found where incident state was SC, NY, OH, NC and PA


```
In [29]: plt.figure(figsize=(15,7))
sns.countplot(x='authorities_contacted',hue='fraud_reported',data=df)
```


```
Out[29]: <AxesSubplot:xlabel='authorities_contacted', ylabel='count'>
```



- Most of the fraud founded where cx have contacted authorities_contacted of Fire, Other and Ambulance

```
In [30]: plt.figure(figsize=(15,7))
sns.countplot(x='incident_severity',hue='fraud_reported',data=df)
```


```
Out[30]: <AxesSubplot:xlabel='incident_severity', ylabel='count'>
```



- Most of fraud reported claim are Major Damage , most of them are counted as fraud
- There are very fraud report in Trivial damage.

```
In [31]: plt.figure(figsize=(15,7))
sns.countplot(x='collision_type',hue='fraud_reported',data=df)
```


```
Out[31]: <AxesSubplot:xlabel='collision_type', ylabel='count'>
```



- Most of fraud reported claims in Side Collision, Rear Collision and Front Collision.
- Some of reported claim we have those are in no info of collision type

```
In [32]: plt.figure(figsize=(15,7))
sns.countplot(x='incident_type',hue='fraud_reported',data=df)
```


```
Out[32]: <AxesSubplot:xlabel='incident_type', ylabel='count'>
```



- We have most claim reported and fraud reported in Single vehicle collision and multi-vehicle collision incident type.
- In other two incident type vehicle theft and parked car are very less fraud report.

```
In [33]: plt.figure(figsize=(15,7))
sns.countplot(x='incident_Month',hue='fraud_reported',data=df)
```


```
Out[33]: <AxesSubplot:xlabel='incident_Month', ylabel='count'>
```



- Most of the data we have reported in month 1 and 2 fraud reported are also around same in both months.

```
In [34]: plt.figure(figsize=(15,7))
sns.violinplot(y='capital-loss',x='fraud_reported',data=df)
```


```
Out[34]: <AxesSubplot:xlabel='fraud_reported', ylabel='capital-loss'>
```



- Most of fraud reported we see in cx those capital loss are and -40000 to -60000

```
In [35]: plt.figure(figsize=(15,7))  
sns.violinplot(y='capital-gains',x='fraud_reported',data=df)
```


```
Out[35]: <AxesSubplot:xlabel='fraud_reported', ylabel='capital-gains'>
```



- Most of the fraud reported are in cx those capital gains are 0 and around 50000

```
In [36]: plt.figure(figsize=(15,7))
sns.countplot(x='insured_relationship',hue='fraud_reported',data=df)
plt.xticks(rotation=70)
```


```
Out[36]: (array([0, 1, 2, 3, 4, 5]),
 [Text(0, 0, 'husband'),
  Text(1, 0, 'other-relative'),
  Text(2, 0, 'own-child'),
  Text(3, 0, 'unmarried'),
  Text(4, 0, 'wife'),
  Text(5, 0, 'not-in-family')])
```



- According to the data info, cx those having insured relationship with other relative and wife are most reported than others.

```
In [37]: plt.figure(figsize=(15,7))
sns.countplot(x='insured_hobbies',hue='fraud_reported',data=df)
plt.xticks(rotation=70)
```

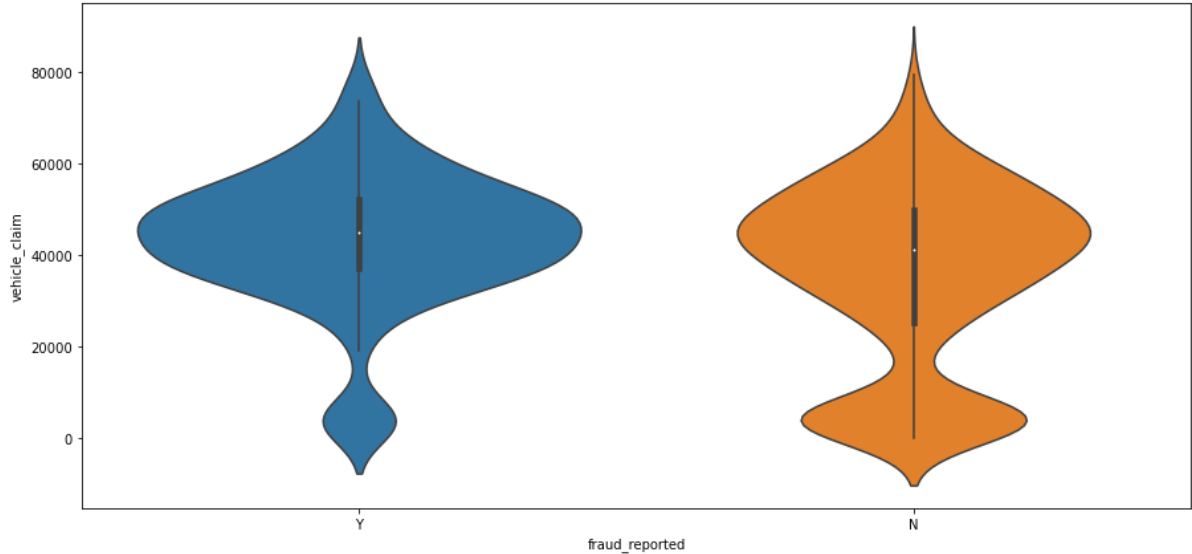
```
Out[37]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19]),
 [Text(0, 0, 'sleeping'),
  Text(1, 0, 'reading'),
  Text(2, 0, 'board-games'),
  Text(3, 0, 'bunge-jumping'),
  Text(4, 0, 'base-jumping'),
  Text(5, 0, 'golf'),
  Text(6, 0, 'camping'),
  Text(7, 0, 'dancing'),
  Text(8, 0, 'skydiving'),
  Text(9, 0, 'movies'),
  Text(10, 0, 'hiking'),
  Text(11, 0, 'yachting'),
  Text(12, 0, 'paintball'),
  Text(13, 0, 'chess'),
  Text(14, 0, 'kayaking'),
  Text(15, 0, 'polo'),
  Text(16, 0, 'basketball'),
  Text(17, 0, 'video-games'),
  Text(18, 0, 'cross-fit'),
  Text(19, 0, 'exercise')])
```



- Here we can see the cx those hobbies are Chess, are most fraud reported cx.
- Here we can see the cx those hobbies are cross-fit , are also most fraud reported cx
- After that cx those hobbies are reading, board games, base-jumping, yechting, painball, polo and etc also most fraud report cx than others

```
In [38]: plt.figure(figsize=(15,7))  
sns.violinplot(y='vehicle_claim',x='fraud_reported',data=df)
```


```
Out[38]: <AxesSubplot:xlabel='fraud_reported', ylabel='vehicle_claim'>
```



- In vehicle_claim most of the cx are between 100 to 6500
- Fraud report is higher in cx are between 3500 to 5500 vehicle_claim

```
In [39]: plt.figure(figsize=(15,7))  
sns.violinplot(y='property_claim',x='fraud_reported',data=df)
```


```
Out[39]: <AxesSubplot:xlabel='fraud_reported', ylabel='property_claim'>
```



- In Property claim fraud reported is higher in cx claimed of 5000 to 8000

```
In [40]: plt.figure(figsize=(15,7))
sns.violinplot(y='injury_claim',x='fraud_reported',data=df)
```


```
Out[40]: <AxesSubplot:xlabel='fraud_reported', ylabel='injury_claim'>
```



- Most of the cx injury claim is 100 to 15000
- Most fraud reported cx are between 5000 to 8000 injury claim

```
In [41]: plt.figure(figsize=(15,7))
sns.countplot(x='insured_education_level', hue='fraud_reported',data=df)
```


```
Out[41]: <AxesSubplot:xlabel='insured_education_level', ylabel='count'>
```



- We higher cx those are education level is High School and least cx in PhD and College
- There is very little difference in fraud report in all kind of cx
- Fraud report is little higher in cx those education level is JD, MD, PhD, and College than others.


```
In [42]: plt.figure(figsize=(15,7))
sns.countplot(x='insured_sex', hue='fraud_reported', data=df)
```


```
Out[42]: <AxesSubplot:xlabel='insured_sex', ylabel='count'>
```



- We have more cx in Female than the Male
- But fraud reported is little high Male cx

```
In [43]: plt.figure(figsize=(15,7))
sns.violinplot(x='fraud_reported', y='umbrella_limit', data=df)
```


```
Out[43]: <AxesSubplot:xlabel='fraud_reported', ylabel='umbrella_limit'>
```



- Most of cx we have their umbrella limit is 0
- fraud reported is also high those cx


```
In [44]: plt.figure(figsize=(15,7))
sns.violinplot(x='fraud_reported',y='policy_annual_premium',data=df)
```

```
Out[44]: <AxesSubplot:xlabel='fraud_reported', ylabel='policy_annual_premium'>
```



- Most of cx we have 1000 to 1500 of premiums payers
- Most of the fraud reported we found in those premium is 1250 to 1300


```
In [45]: plt.figure(figsize=(15,7))
ax=sns.countplot(x='policy_deductable',hue='fraud_reported',data=df)
for p in ax.patches:
    ax.annotate(int(p.get_height()), (p.get_x()+0.25, p.get_height()+1), va='bottom',
                color='black')
```



- Here we notice here according to count fraud report high in cx of 2000 policy_deductable

```
In [46]: plt.figure(figsize=(15,7))
sns.countplot(x='policy_csl',hue='fraud_reported',data=df)
```


```
Out[46]: <AxesSubplot:xlabel='policy_csl', ylabel='count'>
```



- In all three csl we same kind of similarity fraud report is also common in all policy_csl

```
In [47]: plt.figure(figsize=(15,7))
sns.countplot(x='policy_state',hue='fraud_reported',data=df)
```


```
Out[47]: <AxesSubplot:xlabel='policy_state', ylabel='count'>
```



- We have cx from three stats and fraud report almost common in all three states.

```
In [48]: plt.figure(figsize=(15,7))
sns.countplot(x='age',hue='fraud_reported',data=df)
```


```
Out[48]: <AxesSubplot:xlabel='age', ylabel='count'>
```



- We see most of the cx are age of 26 to 50
- Fraud reported is higher in cx 26 to 50 of age.

```
In [49]: plt.figure(figsize=(10,7))  
sns.set(color_codes=True)  
sns.violinplot(y='months_as_customer',x='fraud_reported',data=df)
```


```
Out[49]: <AxesSubplot:xlabel='fraud_reported', ylabel='months_as_customer'>
```



- Most of the cx are 100 to 300 months old
- We can notice here fraud reort is high in cx between 100 to 250 months old

```
In [50]: plt.figure(figsize=(15,7))
sns.countplot(x='policy_bind_Year',hue='fraud_reported',data=df)
plt.xticks(rotation =70,)
```


```
Out[50]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25]),
 [Text(0, 0, '2014'),
  Text(1, 0, '2006'),
  Text(2, 0, '2000'),
  Text(3, 0, '1990'),
  Text(4, 0, '1997'),
  Text(5, 0, '2011'),
  Text(6, 0, '2002'),
  Text(7, 0, '1999'),
  Text(8, 0, '2012'),
  Text(9, 0, '1998'),
  Text(10, 0, '1992'),
  Text(11, 0, '2005'),
  Text(12, 0, '2004'),
  Text(13, 0, '1994'),
  Text(14, 0, '1991'),
  Text(15, 0, '1996'),
  Text(16, 0, '2013'),
  Text(17, 0, '2008'),
  Text(18, 0, '2003'),
  Text(19, 0, '2010'),
  Text(20, 0, '1993'),
  Text(21, 0, '1995'),
  Text(22, 0, '2007'),
  Text(23, 0, '2001'),
  Text(24, 0, '2009'),
  Text(25, 0, '2015')])
```



- Here we notice fraud report is high in policy_vind_Year of 2014, 2006, 1990, 1997, 2002, 1999, 2012, 1998, 1994, 2008, 2003 and etc thanothers

```
In [51]: plt.figure(figsize=(15,7))
sns.countplot(x='auto_year',hue='fraud_reported',data=df)
plt.xticks(rotation =70,)
```


```
Out[51]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20]),
 [Text(0, 0, '1995'),
  Text(1, 0, '1996'),
  Text(2, 0, '1997'),
  Text(3, 0, '1998'),
  Text(4, 0, '1999'),
  Text(5, 0, '2000'),
  Text(6, 0, '2001'),
  Text(7, 0, '2002'),
  Text(8, 0, '2003'),
  Text(9, 0, '2004'),
  Text(10, 0, '2005'),
  Text(11, 0, '2006'),
  Text(12, 0, '2007'),
  Text(13, 0, '2008'),
  Text(14, 0, '2009'),
  Text(15, 0, '2010'),
  Text(16, 0, '2011'),
  Text(17, 0, '2012'),
  Text(18, 0, '2013'),
  Text(19, 0, '2014'),
  Text(20, 0, '2015')])
```



- We can notice we have less vehicle from 1998 and 20010 auto year
- According to vehicle count fraud reported is higher in vehicle of 1996, 2004, 2007, 2011, 2013, 2014 and etc auto year than others

```
In [52]: plt.figure(figsize=(15,7))
sns.countplot(x='auto_model',hue='fraud_reported',data=df)
plt.xticks(rotation =70,)
```


```
Out[52]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                  17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
                  34, 35, 36, 37, 38]),
          [Text(0, 0, '92x'),
           Text(1, 0, 'E400'),
           Text(2, 0, 'RAM'),
           Text(3, 0, 'Tahoe'),
           Text(4, 0, 'RSX'),
           Text(5, 0, '95'),
           Text(6, 0, 'Pathfinder'),
           Text(7, 0, 'A5'),
           Text(8, 0, 'Camry'),
           Text(9, 0, 'F150'),
           Text(10, 0, 'A3'),
           Text(11, 0, 'Highlander'),
           Text(12, 0, 'Neon'),
           Text(13, 0, 'MDX'),
           Text(14, 0, 'Maxima'),
           Text(15, 0, 'Legacy'),
           Text(16, 0, 'TL'),
           Text(17, 0, 'Impreza'),
           Text(18, 0, 'Forrester'),
           Text(19, 0, 'Escape'),
           Text(20, 0, 'Corolla'),
           Text(21, 0, '3 Series'),
           Text(22, 0, 'C300'),
           Text(23, 0, 'Wrangler'),
           Text(24, 0, 'M5'),
           Text(25, 0, 'X5'),
           Text(26, 0, 'Civic'),
           Text(27, 0, 'Passat'),
           Text(28, 0, 'Silverado'),
           Text(29, 0, 'CRV'),
           Text(30, 0, '93'),
           Text(31, 0, 'Accord'),
           Text(32, 0, 'X6'),
           Text(33, 0, 'Malibu'),
           Text(34, 0, 'Fusion'),
           Text(35, 0, 'Jetta'),
           Text(36, 0, 'ML350'),
           Text(37, 0, 'Ultima'),
           Text(38, 0, 'Grand Cherokee']])
```

- Most of auto_model we see in Wrangler, RAM , Pathfinder, Neon and etc.
- We can notice Fraud report is higher in 92x, RAM, E400, Tahore, A5, F150, X5, C300, M5, CIVIC, SILVERADOX6 ML300 and etc auto_models

```
In [53]: plt.figure(figsize=(15,7))
sns.countplot(x='auto_make',hue='fraud_reported',data=df)
```

```
Out[53]: <AxesSubplot:xlabel='auto_make', ylabel='count'>
```



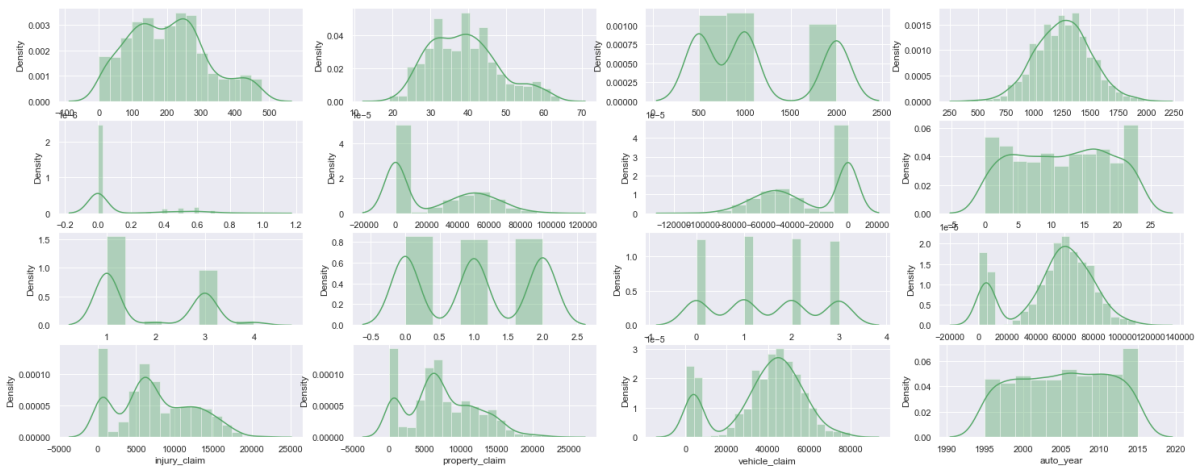
- Fraud reported claim is higher in Saab, Merceded, Dodge, Chevrolet, Audi , Ford, Suburu BMW and volkswagen auto_make
- In Other Auto_make brand is less than other.

Have done the analysis of all important features, some of features don't have much information to define

Skewness Handling

Skew and Outliers will be handel in numerical columns only

```
In [54]: # plotting for numerical columns only
plt.figure(figsize=(25,20))
for i in enumerate(df.select_dtypes(include=['int64','float','int32'])):
    plt.subplot(8,4,i[0]+1)
    sns.distplot(df[i[1]],color='g')
```



```
In [55]: df.select_dtypes(include=['int64','float','int32']).skew()
```

```
Out[55]: months_as_customer    0.364014
age                            0.479796
policy_deductable              0.476426
policy_annual_premium          0.005374
umbrella_limit                 1.806100
capital-gains                  0.477220
capital-loss                   -0.389813
incident_hour_of_the_day       -0.034990
number_of_vehicles_involved    0.501009
bodily_injuries                0.012940
witnesses                      0.018399
total_claim_amount             -0.595646
injury_claim                   0.265382
property_claim                 0.378121
vehicle_claim                  -0.622627
auto_year                     -0.049502
dtype: float64
```

Skewness more than +/-0.5 will we treated

Object and target variable columns will not be treated

we can see here most of the columns are skewed

will only deal with numerical columns

- umbrella_limit
- total_claim_amount
- vehicle_claim

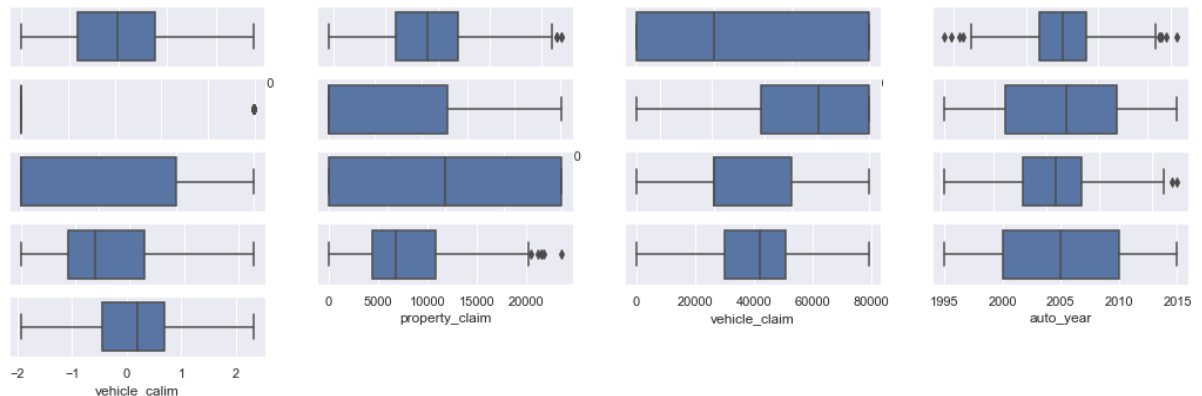
will be treated

```
In [56]: from sklearn.preprocessing import power_transform

df[['umbrella_limit', 'total_claim_amount', 'vehicle_calim']] = power_transform(
    df[['umbrella_limit', 'total_claim_amount', 'vehicle_claim']], method='yeo-jc
```

Outliers Handling

```
In [57]: plt.figure(figsize=(18,10))
for i in enumerate(df.select_dtypes(include=['int64', 'float', 'int32'])):
    plt.subplot(9,4,i[0]+1)
    sns.boxplot(df[i[1]])
```



- Only some of columns seems having outliers after skewness removed

Outliers Removal

ZSCORE Method

```
In [58]: from scipy.stats import zscore

z=np.abs(zscore(df.select_dtypes(include=['int64','float','int32'])))

print(np.where(z>3))

(array([229, 248, 499, 762, 806], dtype=int64), array([ 3,  3, 13,  3,  6], d
type=int64))
```

```
In [59]: df_1=df[(z<3).all(axis=1)]
print(("with outliers::",df.shape))
print("After removing outliers::",df_1.shape)

('with outliers::', (999, 41))
After removing outliers:: (994, 41)
```

- After using zscore method we only lose 5 rows from data

IQR Method

```
In [60]: from scipy import stats
IQR = stats.iqr(df.select_dtypes(include=['int64','float','int32']))
IQR
```

Out[60]: 1997.574235635032

```
In [61]: Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)

df_out = df.select_dtypes(include=['int64','float','int32'])[~((df.select_dtype
print(df_out.shape)

(311, 17)
```

choosing Zscore because there is huge dataloss in IQR

```
In [62]: df= df_1
```

Using LabelEncoder for convering categorical to numerical

```
In [63]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 994 entries, 0 to 999
Data columns (total 41 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   months_as_customer                          994 non-null    int64
1   age                                           994 non-null    int64
2   policy_state                                994 non-null    object
3   policy_csl                                  994 non-null    object
4   policy_deductable                           994 non-null    int64
5   policy_annual_premium                       994 non-null    float64
6   umbrella_limit                             994 non-null    float64
7   insured_sex                                 994 non-null    object
8   insured_education_level                     994 non-null    object
9   insured_occupation                          994 non-null    object
10  insured_hobbies                             994 non-null    object
11  insured_relationship                         994 non-null    object
12  capital-gains                              994 non-null    int64
13  capital-loss                               994 non-null    int64
14  incident_date                               994 non-null    object
15  incident_type                               994 non-null    object
16  collision_type                              994 non-null    object
17  incident_severity                           994 non-null    object
18  authorities_contacted                       994 non-null    object
19  incident_state                              994 non-null    object
20  incident_city                               994 non-null    object
21  incident_hour_of_the_day                    994 non-null    int64
22  number_of_vehicles_involved                 994 non-null    int64
23  property_damage                             994 non-null    object
24  bodily_injuries                             994 non-null    int64
25  witnesses                                   994 non-null    int64
26  police_report_available                     994 non-null    object
27  total_claim_amount                          994 non-null    float64
28  injury_claim                                994 non-null    int64
29  property_claim                              994 non-null    int64
30  vehicle_claim                               994 non-null    int64
31  auto_make                                   994 non-null    object
32  auto_model                                  994 non-null    object
33  auto_year                                   994 non-null    int64
34  fraud_reported                             994 non-null    object
35  incident_Date                               994 non-null    object
36  incident_Month                             994 non-null    object
37  policy_bind_Date                           994 non-null    object
38  policy_bind_Month                          994 non-null    object
39  policy_bind_Year                           994 non-null    object
40  vehicle_calim                               994 non-null    float64
dtypes: float64(4), int64(13), object(24)
memory usage: 326.2+ KB
```

```
In [64]: from sklearn.preprocessing import LabelEncoder
```

```
In [65]: # encoding object columns into Numeric values in df
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
categ_data = df.select_dtypes(exclude=['int64','float','int32'])
for val in categ_data:
    df[val]=le.fit_transform(df[val].astype(str))
```

```
In [66]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 994 entries, 0 to 999
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   months_as_customer                    994 non-null    int64
1   age                                   994 non-null    int64
2   policy_state                          994 non-null    int32
3   policy_csl                           994 non-null    int32
4   policy_deductable                    994 non-null    int64
5   policy_annual_premium                 994 non-null    float64
6   umbrella_limit                       994 non-null    float64
7   insured_sex                          994 non-null    int32
8   insured_education_level              994 non-null    int32
9   insured_occupation                   994 non-null    int32
10  insured_hobbies                       994 non-null    int32
11  insured_relationship                 994 non-null    int32
12  capital-gains                        994 non-null    int64
13  capital-loss                         994 non-null    int64
14  incident_date                        994 non-null    int32
15  incident_type                        994 non-null    int32
16  collision_type                       994 non-null    int32
17  incident_severity                    994 non-null    int32
18  authorities_contacted                994 non-null    int32
19  incident_state                       994 non-null    int32
20  incident_city                        994 non-null    int32
21  incident_hour_of_the_day             994 non-null    int64
22  number_of_vehicles_involved          994 non-null    int64
23  property_damage                      994 non-null    int32
24  bodily_injuries                      994 non-null    int64
25  witnesses                            994 non-null    int64
26  police_report_available              994 non-null    int32
27  total_claim_amount                   994 non-null    float64
28  injury_claim                         994 non-null    int64
29  property_claim                       994 non-null    int64
30  vehicle_claim                        994 non-null    int64
31  auto_make                            994 non-null    int32
32  auto_model                           994 non-null    int32
33  auto_year                            994 non-null    int64
34  fraud_reported                       994 non-null    int32
35  incident_Date                        994 non-null    int32
36  incident_Month                       994 non-null    int32
37  policy_bind_Date                     994 non-null    int32
38  policy_bind_Month                    994 non-null    int32
39  policy_bind_Year                     994 non-null    int32
40  vehicle_calim                        994 non-null    float64
dtypes: float64(4), int32(24), int64(13)
memory usage: 233.0 KB
```

- All columns are converted into numerical now

Dividing data into X and Y

```
In [67]: x = df.drop(['fraud_reported'],axis=1)
         y = df['fraud_reported']
```

```
In [68]: x.shape
```

```
Out[68]: (994, 40)
```

```
In [69]: y.shape
```

```
Out[69]: (994,)
```

Here are the dimension of x and y

Scaling X values

```
In [70]: from sklearn.preprocessing import MinMaxScaler

         sc=MinMaxScaler()
         x=sc.fit_transform(x)
```



```
In [71]: pd.DataFrame(x).isnull().sum()
```

```
Out[71]: 0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
11     0
12     0
13     0
14     0
15     0
16     0
17     0
18     0
19     0
20     0
21     0
22     0
23     0
24     0
25     0
26     0
27     0
28     0
29     0
30     0
31     0
32     0
33     0
34     0
35     0
36     0
37     0
38     0
39     0
dtype: int64
```

So here we can see there is no null value present in the dataset.

```
In [72]: pd.DataFrame(x).describe()
```

```
Out[72]:
```

	0	1	2	3	4	5	6	
count	994.000000	994.000000	994.000000	994.000000	994.000000	994.000000	994.000000	994.0
mean	0.426698	0.443975	0.507042	0.474849	0.424883	0.502476	0.201896	0.4
std	0.240313	0.203319	0.415523	0.402499	0.408046	0.168078	0.401223	0.4
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.0
25%	0.242171	0.288889	0.000000	0.000000	0.000000	0.385783	0.000000	0.0
50%	0.417537	0.422222	0.500000	0.500000	0.333333	0.502305	0.000000	0.0
75%	0.577766	0.572222	1.000000	1.000000	1.000000	0.613049	0.000000	1.0
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.0

8 rows × 40 columns



Here we can see the data have been scaled.


Imbalanced learn

Using Oversampling BoarderlineSMOTE

- because there is less data in Churn for yes

```
In [73]: sns.set_style("whitegrid")
sns.countplot(x='fraud_reported',data=df)
```

```
Out[73]: <AxesSubplot:xlabel='fraud_reported', ylabel='count'>
```




```
In [74]: from imblearn.over_sampling import BorderlineSMOTE
```

```
x_rus, y_rus = BorderlineSMOTE().fit_resample(x,y)
print('Original Target dataset shape:',y.shape)
print('Resample Target dataset shape',y_rus.shape)
```

```
Original Target dataset shape: (994,)
Resample Target dataset shape (1496,)
```

```
In [75]: sns.set_style("whitegrid")
sns.countplot(y_rus)
```

```
Out[75]: <AxesSubplot:xlabel='fraud_reported', ylabel='count'>
```



- Now we have balanced data for model training.

Splitting Train and Test data

```
In [76]: from sklearn.model_selection import train_test_split
```

```
In [77]: x_train,x_test,y_train,y_test = train_test_split(x_rus,y_rus,test_size=.27,rand
```

```
In [78]: x_train.shape
```

```
Out[78]: (1092, 40)
```

```
In [79]: y_train.shape
```

```
Out[79]: (1092,)
```

```
In [80]: x_test.shape
```

```
Out[80]: (404, 40)
```

```
In [81]: y_test.shape
```

```
Out[81]: (404,)
```

So here we successfully split train and test data, now move on for model building

Model Building

```
In [82]: # importing necessary libraries
```

```
from sklearn.linear_model import LogisticRegression

from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import GradientBoostingClassifier

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

import warnings
warnings.filterwarnings('ignore')
```

Logistic Regression

```
In [83]: logreg = LogisticRegression()
logreg_score=cross_val_score(logreg,x_rus,y_rus,cv=5,scoring='accuracy')
print("Cross validation score for logreg:",np.mean(logreg_score))
```

Cross validation score for logreg: 0.7446688963210703

```
In [84]: logreg.fit(x_train,y_train)
LR_predicted=logreg.predict(x_test)

print(accuracy_score(y_test,LR_predicted))
print(confusion_matrix(y_test,LR_predicted))
print(classification_report(y_test,LR_predicted))

print("Training accuracy::",logreg.score(x_train,y_train))
print("Test accuracy::",logreg.score(x_test,y_test))
```

0.7128712871287128

[[122 64]

[52 166]]

	precision	recall	f1-score	support
0	0.70	0.66	0.68	186
1	0.72	0.76	0.74	218
accuracy			0.71	404
macro avg	0.71	0.71	0.71	404
weighted avg	0.71	0.71	0.71	404

Training accuracy:: 0.7628205128205128

Test accuracy:: 0.7128712871287128

Decision Tree Classifier

```
In [85]: dtc = DecisionTreeClassifier()
dtc_score=cross_val_score(dtc,x_rus,y_rus,cv=5,scoring='accuracy')
print("Cross validation score for dtc:",np.mean(dtc_score))
```

Cross validation score for dtc: 0.8369654403567447

```
In [86]: dtc.fit(x_train,y_train)
predicted_dtc=dtc.predict(x_test)

print(accuracy_score(y_test,predicted_dtc))
print(confusion_matrix(y_test,predicted_dtc))
print(classification_report(y_test,predicted_dtc))

print("Training accuracy::",dtc.score(x_train,y_train))
print("Test accuracy::",dtc.score(x_test,y_test))
```

0.806930693069307

[[141 45]

[33 185]]

	precision	recall	f1-score	support
0	0.81	0.76	0.78	186
1	0.80	0.85	0.83	218
accuracy			0.81	404
macro avg	0.81	0.80	0.80	404
weighted avg	0.81	0.81	0.81	404

Training accuracy:: 1.0

Test accuracy:: 0.806930693069307

KNeighbors Classifier

```
In [87]: knn = KNeighborsClassifier()
knn_score=cross_val_score(knn,x_rus,y_rus,cv=5,scoring='accuracy')
print("Cross validation score for knn :",np.mean(knn_score))
```

Cross validation score for knn : 0.6597703455964326

```
In [88]: knn.fit(x_train,y_train)
predicted_knn=knn.predict(x_test)

print(accuracy_score(y_test,predicted_knn))
print(confusion_matrix(y_test,predicted_knn))
print(classification_report(y_test,predicted_knn))

print("Training accuracy::",knn.score(x_train,y_train))
print("Test accuracy::",knn.score(x_test,y_test))
```

0.6633663366336634

[[61 125]

[11 207]]

	precision	recall	f1-score	support
0	0.85	0.33	0.47	186
1	0.62	0.95	0.75	218
accuracy			0.66	404
macro avg	0.74	0.64	0.61	404
weighted avg	0.73	0.66	0.62	404

Training accuracy:: 0.7261904761904762

Test accuracy:: 0.6633663366336634

Random Forest Classifier

```
In [89]: rfc = RandomForestClassifier()
rfc_score=cross_val_score(rfc,x_rus,y_rus,cv=5,scoring='accuracy')
print("Cross validation score for rfc :",np.mean(rfc_score))
```

Cross validation score for rfc : 0.8744147157190636

```
In [90]: rfc.fit(x_train,y_train)
predicted_rfc=rfc.predict(x_test)

print(accuracy_score(y_test,predicted_rfc))
print(confusion_matrix(y_test,predicted_rfc))
print(classification_report(y_test,predicted_rfc))

print("Training accuracy::",rfc.score(x_train,y_train))
print("Test accuracy::",rfc.score(x_test,y_test))
```

0.844059405940594

[[157 29]

[34 184]]

	precision	recall	f1-score	support
0	0.82	0.84	0.83	186
1	0.86	0.84	0.85	218
accuracy			0.84	404
macro avg	0.84	0.84	0.84	404
weighted avg	0.84	0.84	0.84	404

Training accuracy:: 1.0

Test accuracy:: 0.844059405940594

Ensemble Technique

1. AdaBoost Classifier

```
In [91]: adb = AdaBoostClassifier()
adb_score=cross_val_score(adb,x_rus,y_rus,cv=5,scoring='accuracy')
print("Cross validation score for adb :",np.mean(adb_score))
```

Cross validation score for adb : 0.8383389074693424


```
In [92]: adb.fit(x_train,y_train)
predicted_adb=adb.predict(x_test)

print(accuracy_score(y_test,predicted_adb))
print(confusion_matrix(y_test,predicted_adb))
print(classification_report(y_test,predicted_adb))

print("Training accuracy::",adb.score(x_train,y_train))
print("Test accuracy::",adb.score(x_test,y_test))
```

0.8193069306930693

[[152 34]

[39 179]]

	precision	recall	f1-score	support
0	0.80	0.82	0.81	186
1	0.84	0.82	0.83	218
accuracy			0.82	404
macro avg	0.82	0.82	0.82	404
weighted avg	0.82	0.82	0.82	404

Training accuracy:: 0.8946886446886447

Test accuracy:: 0.8193069306930693

2. Bagging Classifier

```
In [93]: bgc = BaggingClassifier()
bgc_score=cross_val_score(bgc,x_rus,y_rus,cv=5,scoring='accuracy')
print("Cross validation score for bgc :",np.mean(bgc_score))
```

Cross validation score for bgc : 0.8730479375696767

```
In [94]: bgc.fit(x_train,y_train)
         predicted_bgc=bgc.predict(x_test)

         print(accuracy_score(y_test,predicted_bgc))
         print(confusion_matrix(y_test,predicted_bgc))
         print(classification_report(y_test,predicted_bgc))

         print("Training accuracy::",bgc.score(x_train,y_train))
         print("Test accuracy::",bgc.score(x_test,y_test))
```

0.8564356435643564

[[155 31]

[27 191]]

	precision	recall	f1-score	support
0	0.85	0.83	0.84	186
1	0.86	0.88	0.87	218
accuracy			0.86	404
macro avg	0.86	0.85	0.86	404
weighted avg	0.86	0.86	0.86	404

Training accuracy:: 0.9935897435897436

Test accuracy:: 0.8564356435643564

3. Gradient Boosting Classifier

```
In [95]: grbc = GradientBoostingClassifier()
         grbc_score=cross_val_score(grbc,x_rus,y_rus,cv=5,scoring='accuracy')
         print("Cross validation score for grbc :",np.mean(grbc_score))
```

Cross validation score for grbc : 0.8790858416945374

```
In [96]: grbc.fit(x_train,y_train)
predicted_grbc=grbc.predict(x_test)

print(accuracy_score(y_test,predicted_grbc))
print(confusion_matrix(y_test,predicted_grbc))
print(classification_report(y_test,predicted_grbc))

print("Training accuracy::",grbc.score(x_train,y_train))
print("Test accuracy::",grbc.score(x_test,y_test))
```

```
0.8910891089108911
```

```
[[156  30]
```

```
 [ 14 204]]
```

	precision	recall	f1-score	support
0	0.92	0.84	0.88	186
1	0.87	0.94	0.90	218
accuracy			0.89	404
macro avg	0.89	0.89	0.89	404
weighted avg	0.89	0.89	0.89	404

```
Training accuracy:: 0.9734432234432234
```

```
Test accuracy:: 0.8910891089108911
```

Observation :

Choosing :-

- Gradient Boosting Classifier as final model for Hyper Parameter Tuning because both train and test accuracies are close and highest as well
- Rest of the models having huge difference between train and test accuracies so not considering them.

Hyper Parameter Tuning : GradientBoostingClassifier

```
In [97]: adb=GradientBoostingClassifier()
param_grid={
    'criterion' :['mse','mae'],
    'n_estimators' :[100,200],
    'learning_rate' :[0.1,0.5,1.0],
    'random_state' :[5],
}
```

```
In [98]: adb_grid=GridSearchCV(GradientBoostingClassifier(),param_grid,cv=4,scoring='acc
```

```
In [99]: adb_grid.fit(x_train,y_train)
adb_pred=adb_grid.best_estimator_.predict(x_test)
print("Accuracy after parameter tuning::",accuracy_score(y_test,adb_pred))
```

Fitting 4 folds for each of 12 candidates, totalling 48 fits
Accuracy after parameter tuning:: 0.8589108910891089

```
In [100]: adb_grid.best_params_
```

```
Out[100]: {'criterion': 'mse',
           'learning_rate': 1.0,
           'n_estimators': 200,
           'random_state': 5}
```

Model Training with best parameters

```
In [101]: best_param={
           'criterion' :['mse'],
           'n_estimators' :[200],
           'learning_rate' :[0.1],
           'random_state' :[5],
           }
```

```
In [102]: best_adb_grid=GridSearchCV(GradientBoostingClassifier(),best_param,cv=4,scoring
```

```
In [103]: best_adb_grid.fit(x_train,y_train)
best_adb_pred=best_adb_grid.best_estimator_.predict(x_test)
print("Accuracy after parameter tuning::",accuracy_score(y_test,best_adb_pred))
```

Fitting 4 folds for each of 1 candidates, totalling 4 fits
Accuracy after parameter tuning:: 0.8910891089108911

As we notice after Hyper Parameter Tuning models accuracy score got increased

Report of GradientBoostingClassifier

```
In [104]: print("Classification Report::\n",classification_report(y_test,best_adb_pred))
```

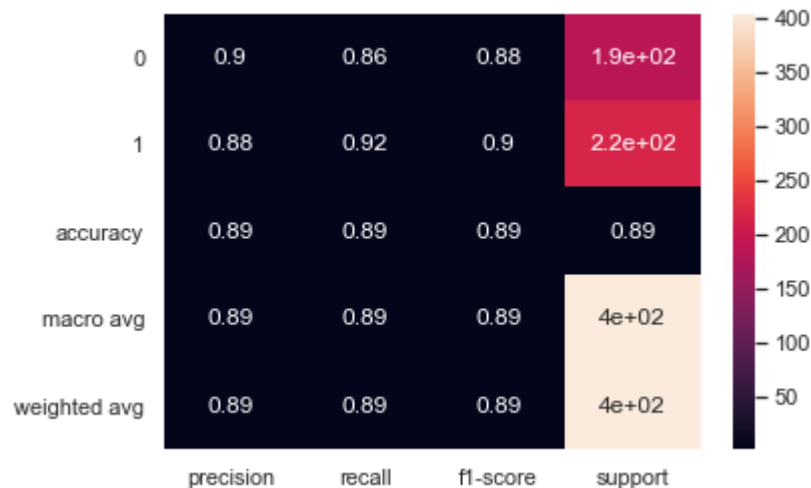
```
Classification Report::
              precision    recall  f1-score   support

     0       0.90      0.86      0.88       186
     1       0.88      0.92      0.90       218

 accuracy          0.89
 macro avg         0.89      0.89      0.89
 weighted avg      0.89      0.89      0.89
```

```
In [105]: clsf_repo = classification_report(y_test,best_adb_pred,output_dict=True)
sns.heatmap(pd.DataFrame(clsf_repo).T, annot=True)
```

```
Out[105]: <AxesSubplot:>
```



Plotting AUC ROC curve

```
In [106]: from sklearn.metrics import roc_auc_score
```

```
In [107]: print("roc auc score::",roc_auc_score(y_test,best_adb_pred))
```

```
roc auc score:: 0.8888231232119956
```


```
In [108]: from sklearn.metrics import roc_curve
```

```
In [109]: fpr1, tpr1, thresh1 = roc_curve(y_test, best_adb_pred, pos_label=1)
```

```
In [110]: # plotting ROC CURVE
sns.set_theme(style="darkgrid")
plt.plot(fpr1, tpr1, linestyle='--', color='red', label='GradientBoostingclassifier')

plt.title("ROC curve")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='best')
```

Out[110]: <matplotlib.legend.Legend at 0x1667102ef40>



Saving the Model

```
In [111]: import joblib
joblib.dump(best_adb_grid.best_estimator_, 'InsuranceClaims-FraudDetection.obj')
```

Out[111]: ['InsuranceClaims-FraudDetection.obj']

So here save the best model using joblib library.

Prediction Result

```
In [112]: # Loading the saved model
model=joblib.load("InsuranceClaims-FraudDetection.obj")

# Prediction
prediction = model.predict(x_test)
prediction
```

```
Out[112]: array([0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1,
1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1,
0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1,
0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1,
1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0,
0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0,
1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0,
1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1,
0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1,
0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0,
0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1,
1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0,
1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0,
1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1,
1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1,
1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0,
0, 1, 1, 0, 1, 1, 0, 0])
```

```
In [115]: pd.DataFrame([model.predict(x_test)[:],y_test[:]],index=["Predicted","Original"])
```

```
Out[115]:
```

	Predicted	Original
0	0	0
1	1	1
2	1	1
3	0	0
4	0	1
...
399	0	0
400	1	1
401	1	1
402	0	0
403	0	1

404 rows × 2 columns

So here we can observe that the actual predicted values are almost same, that means our model worked well.

